

---

# Image Mosaicing

Sahasrajit Anantharamakrishnan<sup>1</sup>, Eli MacColl<sup>2</sup>

<sup>1</sup>Department of Electrical and Computer Engineering, [anantharamakrishn.sa@northeastern.edu](mailto:anantharamakrishn.sa@northeastern.edu)

<sup>2</sup>Department of Electrical and Computer Engineering, [maccoll.e@northeastern.edu](mailto:maccoll.e@northeastern.edu)

March 17, 2023

---

## Abstract

In this project, we aim to develop a computer vision algorithm for image mosaicing with two images by applying a Harris corner detector, identifying corresponding features, and estimating a homography between them. The homography describes how one image can be warped into the coordinate system of the second image to create a mosaic that has the union of all pixels in the two images where overlapping pixels are blended via a blending scheme. Normalized cross-correlation (NCC) is used to find corresponding features between two images such that only the best matches for each corner feature are considered. The homography is estimated using the random sample consensus (RANSAC) algorithm which iteratively identifies the homography that produces the largest set of inliers amongst the noisy correspondences. A least-squares homography can then be calculated using all of the inliers from that largest set of inliers and then used for warping one image onto the other. Using this process, we can effectively generate an image mosaic using two input images.

---

## 1. Introduction

Image mosaicing is a fundamental problem in computer vision that involves stitching multiple images together to form a larger, panoramic view. It is a powerful technique that has a wide range of applications across a variety of fields such as 3D modeling, surveillance, medical imaging, photography, and many more. The process of image mosaicing typically consists of several steps, including feature detection, feature matching between the images, image alignment using those corresponding features, and finally image warping and blending to create a seamless composite image. There are several options for algorithms and techniques to carry out these steps.

The methods we are implementing include a Harris corner detector for feature detection, compute normalized cross-correlation (NCC) for feature matching and use random sample consensus (RANSAC) for determining the homography for image alignment. The Harris corner detector is a popular algorithm for detecting corners in images. It is used to identify significant features in images that can be used for matching and registration.

NCC is a technique used to identify corresponding features in different images. The NCC scores correspond to how close of a match two features are on a scale of -1 to 1. NCC is a powerful method for feature matching. Homography estimation is used to estimate the geometric relationship between two images. It is a transformation matrix that maps one image onto another image, allowing the two to be combined. In order to robustly estimate the homography, RANSAC is used on the noisy correspondences to determine the homography that produces the largest set of inliers that should then be used to estimate the homography for warping.

The goal of this project is to implement an image mosaicing algorithm that uses a Harris corner detector to find features in the images, compute NCC to identify corresponding features, use RANSAC to eliminate outliers and estimate the homography matrix between the images, and then warp one image into the coordinate system of the other to create a mosaic.

## 2. Algorithms

### 2.1. Harris Corner Detector

Harris corner detector is a feature detection algorithm that is used to identify points of interest in an image. These key points are significant because they are the required inputs for various other tasks. In our case, they are used for feature matching. The Harris corner detector gives a mathematical approach for determining changes in intensity within small sections/windows as they are moved across an image. In particular, it looks for areas where the intensity changes in multiple directions, meaning they are likely to be corners or edges. The algorithm computes a corner response function for each pixel in the image using the image gradients to ultimately determine if a given pixel is a corner.

---

#### Algorithm 1 Harris Corner Detector

---

**Input:** Image ( $I$ ).

**Input:** Empirically determined constant ( $k$ ) in the range  $0.04 \leq x \leq 0.06$ . Defaults to 0.04.

**Input:** Window Size ( $N$ ) of ( $N \times N$ ) used for SOBEL mask and Gaussian Blur.

- 1: Compute the image gradients,  $I_x$  and  $I_y$ , with the SOBEL masks  $G_x^S$  and  $G_y^S$  in the  $x$  and  $y$  direction respectively:

$$I_x = G_x^S \times I \quad (1)$$

$$I_y = G_y^S \times I \quad (2)$$

- 2: Compute products of derivatives at each pixel:

$$I_x^2 = I_x \times I_x \quad (3)$$

$$I_y^2 = I_y \times I_y \quad (4)$$

$$I_x I_y = I_x \times I_y \quad (5)$$

- 3: Compute the sums of the products at each pixel using a window averaging:

$$S_x^2 = G_{S'} \times I_x^2 \quad (6)$$

$$S_y^2 = G_{S'} \times I_y^2 \quad (7)$$

$$S_{xy} = G_{S'} \times I_{xy} \quad (8)$$

▷ Here  $G_{S'}$  is a Gaussian Mask

- 4: Define the Matrix at each pixel:

$$M = \begin{bmatrix} S_x^2 & S_{xy} \\ S_{xy} & S_y^2 \end{bmatrix} \quad (9)$$

- 5: Compute the Response ( $R$ ):

$$R = \det(M) - k \cdot (\text{trace}(M))^2 \quad (10)$$

- 6: Threshold  $R$ .

- 7: Compute Nonmax suppression.
-

The Equation for Harris Corner Detector is:

$$E(u, v) = \sum_{x,y} \underbrace{w(x, y)}_{\text{Window Function}} \left[ \underbrace{I(x+u, y+v) - I(x, y)}_{\text{Change in Intensity}} \right]^2 \quad (11)$$

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix} \quad (12)$$

Where,

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \quad (13)$$

## 2.2. Normalized Cross Correlation

Normalized Cross Correlation (NCC) is a technique used in image processing for measuring the similarity between two image patches. Higher NCC scores coincide with the best match and mean it is less likely that the match was identified incorrectly. The NCC values range from  $-1$  to  $1$  where  $-1$  means completely uncorrelated and  $1$  indicates a perfect match. NCC is a normalized version of cross correlation that mitigates an issue by comparing the same area in a scene with different illumination intensities. This is accomplished by normalizing the pixels in the patches before comparing them. NCC involves subtracting the mean from the pixels in the image patch, dividing each pixel by the standard deviation, and computing the cross-correlation as the sum of the products of the normalized image patches. Each patch is a unit norm vector and the NCC is their dot product. In this project, NCC is used during feature matching to identify corresponding points across the two images.

The Equation of Normalized Cross Correlation is given by:

$$N_{fg} = C_{\hat{f}\hat{g}} = \sum_{[i,j] \in \mathbb{R}} \hat{f}(i, j) \hat{g}(i, j) \quad (14)$$

Where,

$$\hat{f} = \frac{f}{\|f\|}; \quad \hat{g} = \frac{g}{\|g\|} \quad (15)$$

## 2.3. RANdom SAMple Consensus (RANSAC)

RANdom SAMple Consensus, or RANSAC, is a robust estimation algorithm that is used to find the best model that fits a set of noisy data. In this project, RANSAC is used to estimate the homography between two images based on a set of noisy correspondences. The algorithm, for lines, involves drawing a sample from the data, uniformly and at random, fitting to that sample, testing the distance from every other data point to that fitted line in order to determine if it is an inlier or an outlier, and repeating that process a certain number of times to find the best fit.

Parameters such as how many points to sample, the distance threshold, and the number of iterations to perform depend on the task and the data being used. The equations for determining the number of samples,  $N$ , are given below. It is dependent on the desired probability of finding a sample free from outliers. For example, a value of  $p = 0.99$  for a certain  $N$  gives a 99% chance of getting a good sample of all inliers within the corresponding number of iterations. The number of iterations required is determined using Table 1 below based on the number of points per sample size,  $s$ , and the probability of a point being an outlier,  $e$ . For our purposes, we require a sample of  $s = 4$  points and opted to perform 72 iterations with a distance threshold of 5 pixels. This means that for a given iteration, points with a distance from the fitted line less than 5 pixels are considered inliers and those with a distance greater than 5 pixels are outliers. After 72 iterations, the largest set of inliers is used to estimate the homography.

Equation (17) and Table 1 gives the number of samples (Iteration) to compute given other parameters.

$$p = 1 - (1 - (1 - e)^s)^N \quad (16)$$

Rearranging to get the value of  $N$ :

$$N = \frac{\ln(1 - p)}{\ln(1 - (1 - e)^s)} \quad (17)$$

Here,

- $e$  — Probability that a point is an outlier.
- $s$  — Number of points in a sample.
- $N$  — Number of samples (We want to compute this).
- $p$  — Desired probability that we get a good sample.

Proportion of outliers $e$							
$s$	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

**Table 1.** Choosing the number of iterations  $N$  given a sample size  $s$ .

---

**Algorithm 2** RANdom SAMple Consensus (RANSAC)

---

**Input:**

- $n$  — the smallest number of points required.
- $k$  — the number of iterations required.
- $t$  — the threshold used to identify a point that fits well.
- $d$  — the number of nearby points required to assert a model fits well.

**Algorithm:**

```

1: repeat
2:   Draw a sample of  $n$  point from the data uniformly and at random.
3:   Fit to that set of  $n$  points.
4:   for all points outside the sample do
5:     if the distance from the point to the line is less than  $t$  then
6:       The point is close.
7:     else
8:       The point is far.
9:     end if
10:  end for
11:  if there are  $d$  or more points close to the line then
12:    There is a good fit.
13:    Refit the line using all these points.
14:  end if
15: until  $k$  iterations have occurred.
16: Use the best fit from this collection, using the fitting error as a criterion.
```

---

## 2.4. Homography Estimation

A homography is a mathematical transformation that describes the relationship between two coordinate systems. For the purposes of this project, we estimate a homography to map coordinates of points in one image to points in the other image. This is necessary for the final warping step to create a single composite image as a mosaic of the two images. The equations for homography estimation can be seen below.

Equations to find Homography Matrix ( $H$ ):

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \sim \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad (18)$$

In Equation form:

$$x' = \frac{h_{11}x + h_{12}y + h_{13}}{h_{31}x + h_{32}y + h_{33}} \quad (19)$$

$$y' = \frac{h_{21}x + h_{22}y + h_{23}}{h_{31}x + h_{32}y + h_{33}} \quad (20)$$

Since we are using the algebraic distance method to solve the homography matrix the above equation holds true only when the condition  $\|h\| = 1$  is met. So,

$$(h_{31}x + h_{32}y + h_{33}) \cdot x' = h_{11}x + h_{12}y + h_{13} \quad (21)$$

$$(h_{31}x + h_{32}y + h_{33}) \cdot y' = h_{21}x + h_{22}y + h_{23} \quad (22)$$

$$\implies h_{11}x + h_{12}y - h_{31}xx' - h_{32}yx' - h_{33}x' + h_{13} = 0 \quad (23)$$

$$\implies h_{21}x + h_{22}y - h_{31}xy' - h_{32}yy' - h_{33}y' + h_{23} = 0 \quad (24)$$

In Matrix form:

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x'_1 & -y_1x'_1 & -x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1y'_1 & -y_1y'_1 & -y'_1 \\ & & & & & \vdots & & & \end{bmatrix} \cdot \begin{bmatrix} h_{11} \\ h_{12} \\ \vdots \\ h_{33} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \quad (25)$$

Simply,

$$\underbrace{A}_{2N \times 9} \cdot \underbrace{h}_{9 \times 1} = \underbrace{0}_{2N \times 1} \quad (26)$$

The above Homogenous matrix equation can be solved using Singular Value Decomposition (SVD).

## 2.5. Image Warping and Blending

Warping and blending are the final steps in producing a mosaic of multiple images. Prior to warping, the size of the output image must be determined so that it is big enough to contain the union of all pixels in the two images. The image that does not need to be warped is first placed into the output frame. The second is then warped into the output image using the estimated homography or potentially it's inverse. A blending scheme is used in order to combine colors in areas of overlapping pixels. There are several blending schemes, including straight averaging, feathering which considers the distance from the image border, and equalization for adjusting intensities.

Approaches to Blending:

### 1. Straight Averaging

$$P = \frac{(P_1 + P_2)}{2} \quad (27)$$

## 2. Feathering

$$P = \frac{(w_1 * P_1 + w_2 * P_2)}{(w_1 + w_2)} \quad (28)$$

Where  $w_i$  is the distance from the image border

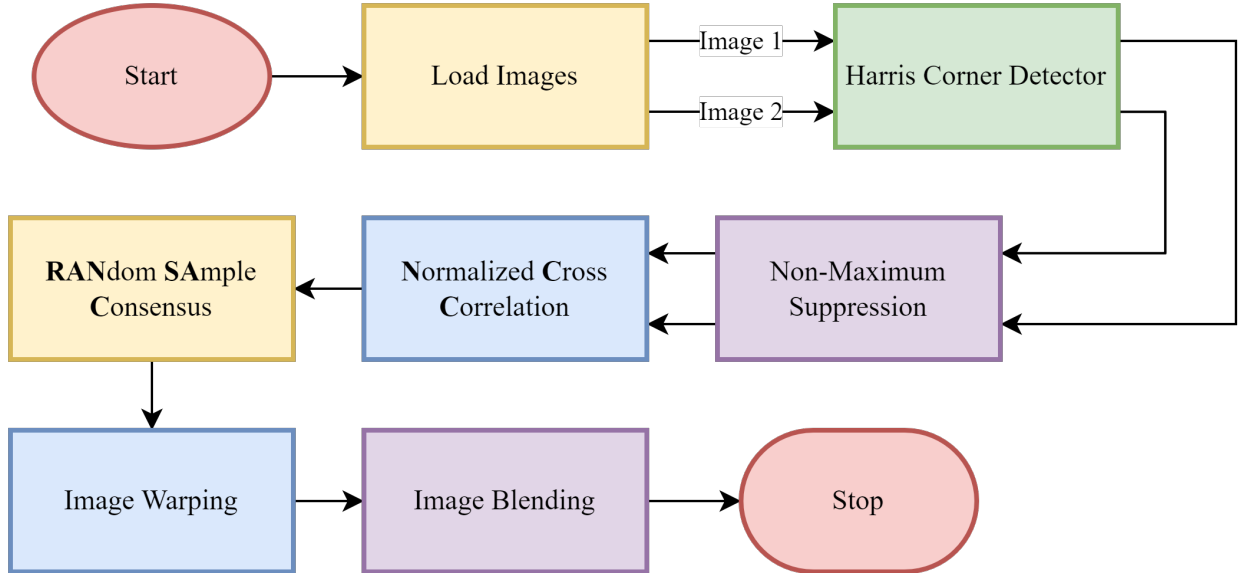
## 3. Equalize intensity statistics (gain, offset)

### 3. Experiments

The workflow of the program goes as follows:

- (a) Read in two image.
- (b) Apply Harris corner detector to each image to identify corner features.
- (c) Find correspondences between the two images by computing NCC to determine the best matches.
- (d) Estimate the homography with the above correspondences using RANSAC for robust estimation.
- (e) Warp one image onto the other, blending overlapping pixels to create a single composite image.

Following this fundamental outline, we implemented the techniques described in the Algorithms section above in order to determine the values of parameters and develop our image mosaicing program.



**Figure 1.** Image Mosaicing Flowchart.

#### 3.1. Hallway Dataset


The hallway dataset contains images of a hallway scene from 3 distinct angles. They contain little visual clutter and contain a few large distinguishable features, such as bulletin board, doors, and a water fountain. Additionally, the illumination intensity appears relatively uniform throughout the scene.



**Figure 2.** Hallway dataset - Input Images

### 3.2. Office Dataset

The office dataset contains images of an office scene from 10 distinct angles, effectively providing a 360° panoramic view of the room across the set of images. This scene is more complex than the hallway scene in the sense that there is a lot more going on, visually. There are numerous small objects scattered across a cluttered desk and some reflective surfaces such as the glass on a picture frame and a DVD. Moreover, there is a window as a main source of light, causing the illumination intensity to differ throughout the scene in places like underneath the desk.



**Figure 3.** Office dataset - Input Images

## 4. Results and Discussion

The images displayed in the following results represent an image pair from both datasets, the hallway, and the office. These select outputs aim to best demonstrate our overall results since we cannot include them all for the sake of brevity.

### 4.1. Feature Detection

Following the program outline described in the Experiments section above, the first step after loading the image pair being used to create a mosaic is to perform feature detection. We apply a Harris corner detector over each image and then do non-maximum suppression to obtain a sparse set of corner features in both images. A Harris corner detector is able to identify corner features by sliding a small window across the image and evaluating changes in intensity.

In our implementation of the Harris corner detector algorithm, we used a  $3 \times 3$  window size to apply both the Sobel and Gaussian masks following the procedure in Algorithm 1. The response,  $R$ , was computed using

an empirically determined constant  $k$  with a value of 0.04, which falls within the range of  $k \in [0.04, 0.06]$ . To reduce the number of detected corner points, we employed the use of Non-Maximum Suppression (NMS) on the computed  $R$  values. Subsequently, we applied a dataset-specific threshold to the suppressed values in order to remove any remaining noise in the response matrix. The threshold values were determined through trial and error to be 0.0025 and 0.01 for the hallway and office dataset respectively.

As depicted in Figures 4 and 5, we were able to accurately detect corner features in the images. That being said, there were some inaccuracies in our results in the form of false positives and false negatives. For instance, in Figures 4a and 4b a false positive was detected near the top right of the wooden door due to the reflection of the overhead lighting. Likewise, in Figures 5a and 5b, a corner was detected on the glass of the framed blue image due to the reflection of light from the window. Additionally, there were false negatives in both datasets. For example, the top left corner of the bulletin board went undetected in Figure 4a, and the top and bottom left corners of the framed blue picture were missed in Figure 5a.



**Figure 4.** Hallway dataset - Feature Detection



**Figure 5.** Office dataset - Feature Detection

## 4.2. Feature Matching

Using the outputs from the feature detection step, feature matching is performed in order to find the correspondences between the two images. Given a set of two corners from the two images, NCC is computed using image patches centered around each corner. Every corner detected in one image is compared to every corner detected in the other image in order to determine the best possible match, making this an  $O(n^2)$  process. We also set a threshold such that the best possible matches have a high NCC score in an effort to eliminate incorrect matches. In our implementation of NCC, we opted for a patch size of 5 pixels.


The results produced by our feature-matching algorithm proved to be sufficient for our purposes. Our NCC function identified enough correspondences to be able to proceed with RANSAC for homography estimation, but not an abundance of extraneous matches that would be considered outliers as seen in Figures 6



and 7. The lines in the figures are drawn between corresponding points in the scene across the two images. High NCC scores indicate correct matches. Corresponding corners that had the best NCC score comparatively are included in the output correspondences to be used in future steps. High NCC values also provides a level of confidence in the quality of the matches being identified. Moreover, our correspondence visualization closely resembles that of the example provided in the project description.



**Figure 6.** Hallway dataset - Feature Matching



**Figure 7.** Office dataset - Feature Matching

### 4.3. Homography Estimation

With the correspondences identified from the feature matching process, we are now able to estimate the homography required to warp one image into the coordinate system of the second. Given that the correspondences we found are likely to have many errors/outliers which is exhibited in figures 5 and 6 by the various slopes of the lines, we perform RANSAC to robustly estimate the homography from the noisy correspondences. For an overview of the process, we use a random sample of 4 correspondences to estimate a homography and then determine the number of inliers produced after mapping all points using that homography. This iterative process is completed a certain number of times before computing a least-squares homography using all of the inliers from the largest set of inliers that was found. The RANSAC algorithm and parameters are described in greater detail in the algorithms section above in Algorithm 2.


In our implementation of RANSAC, we used a sample size of  $N = 4$  points, a distance threshold of 5 pixels, and ran it over 72 iterations as per the highlighted row in Table 1 above. We were able to achieve positive results using these parameters as shown by visualizations in Figures 8 and 9. The lines in the figures are drawn between corresponding features in the scene across the two images. When compared to Figures 6 and 7 which display all of the correspondences that were found including errors and outliers, it is clear that the set of correspondences that our RANSAC determined to be inliers more accurately reflect corresponding features in the two images. This can be seen either through individual visual inspection of the lines, or

more generally, the similar slopes of all the correspondence lines between the two images. This makes sense because the two images capture the same scene with a significant amount of overlapping features but taken at slightly different angles. Therefore, lines drawn between correctly matched features in the two images should all have the same slope.

Once RANSAC is complete, we have the largest set of inliers found across all of the iterations. Those inliers are used to compute a final homography that enables us to warp one image into the coordinate system of the other image. The final homography matrices found for each dataset can be seen below in Equations (29) and (30). These are the homographies used to warp one image onto the other in the final output frame for the corresponding dataset.



**Figure 8.** Hallway dataset - RANSAC Inliers



**Figure 9.** Office dataset - RANSAC Inliers

$$H_{\text{Hallway}} = \begin{bmatrix} 1.16410639 \times 10^0 & 7.24666331 \times 10^{-2} & -2.59175335 \times 10^1 \\ 1.30681317 \times 10^{-2} & 1.25005557 \times 10^{-2} & -1.57003065 \times 10^2 \\ 2.45925764 \times 10^5 & 4.84730417 \times 10^{-4} & 1.00000000 \times 10^0 \end{bmatrix} \quad (29)$$

$$H_{\text{Office}} = \begin{bmatrix} 1.09010727 \times 10^0 & -9.97310911 \times 10^{-2} & 5.51610087 \times 10^0 \\ 1.88851954 \times 10^{-1} & 7.99619183 \times 10^{-1} & 1.35065914 \times 10^2 \\ 9.03209811 \times 10^{-4} & -9.85230262 \times 10^{-4} & 1.00000000 \times 10^0 \end{bmatrix} \quad (30)$$

#### 4.4. Warping and Blending

Using the homography matrix estimated previously with RANSAC, we are now able to map points from the coordinate system of one image to the coordinate system of another. This is required for the final step of warping one image onto the other and blending the overlapping pixels together to create a single image that shows the union of all pixels from both images.

In order to accomplish this, we needed to create a final output image that was large enough to contain the composition of the two images. This is primarily a concern for the width of the output image given the nature of the input images being taken at similar vertical angles. With that in mind, we are effectively creating a panorama, so we determined a sufficient width of the final image to be the sum of the widths of the two input images. We proceeded to insert the image not being warped into the output frame and then warped in the second image. As shown in Section 4.4, we were then able to successfully produce a mosaic of the two inputs images for each dataset.



**Figure 10.** Hallway dataset - Output Mosaic



**Figure 11.** Office dataset - Output Mosaic



#### 4.5. Extra Credit

The extra credit portion of this project involves warping one image into a user-selected region in a second image. To accomplish this, we consider the corner points of the image being embedded to be the corners of the image. To find the corners of the region where the image is being embedded, we take user input in the form of mouse clicks on the target image in clockwise order starting with the top left corner. Using the corner coordinates from both images, we can calculate a homography matrix to map points from one image coordinate system to the other. The computed homography is then used to warp the image we want to embed into the base image, creating a single composite image.

We managed to accomplish this task as shown below in Figure 12 where we successfully embedded an image onto the bulletin board in the hallway scene. We selected this region of the image by using mouse clicks to pick the corners of our desired frame in the order that was previously mentioned above. The locations of the mouse clicks are represented by the red squares centered around the click coordinates. The order of the corner selection is significant here because it influences the orientation of the image by defining the coordinate system of the region we are embedding the image into. For example, if the corners of the target region were selected clockwise starting from the bottom right corner instead, the resulting embedded image would be upside down.



**Figure 12.** Extra Credit - Output Example

#### 5. Conclusion

In conclusion, the project "Image Mosaicing" demonstrates a practical application of image processing techniques to create a mosaic using two images of the same scene. The project defines a general procedure of feature detection, feature matching, homography estimation, and warping/blending. Implementing and performing this process can be used to find corner features in the images to then identify correspondences that are crucial for image alignment to ultimately produce a single composite image depicting the scene in the images.

The performance of the image mosaicing algorithm is dependent on several factors and can influence the decision as to what techniques to use for the various steps. For example, in scenes that vary in illumination

intensity, NCC is a good choice for feature matching because normalization mitigates the impact of the different illumination intensities on the correlation results. Furthermore, the choice of parameters for each technique also plays a significant role. The project provides resources that describe how to calculate and fine-tune the parameters for some of these techniques, such as RANSAC for homography estimation.

Overall, this project demonstrates how fundamental image processing techniques can be powerful tools for image mosaicing. By applying these techniques to a set of images depicting a single scene, it is possible to extract useful information about corresponding points in the images and create a seamless composite image.

## Appendix - Code

Checkout our source code in [GitHub](https://github.com/Sahas-Ananth/CVP2) (<https://github.com/Sahas-Ananth/CVP2>)

### A. Source Code

```
1 import os
2 from random import choices, randint
3
4 import cv2
5 import numpy as np
6
7
8 class Panorama:
9     def __init__(self, path: str) -> None:
10         """Initialize the Panorama class.
11
12         Args:
13             path (str): Path to the folder containing the images.
14         """
15         if path == "DanaHallWay1":
16             self.harris_thresh = 0.0025
17         elif path == "DanaOffice":
18             self.harris_thresh = 0.01
19
20     def load_images(self, path: str) -> np.ndarray:
21         """Load images from a path.
22
23         Args:
24             path (string): Path to the folder containing the images.
25
26         Returns:
27             np.ndarray: Color images.
28             np.ndarray: Grayscale images.
29         """
30         files = [
31             os.path.join(path, f)
32             for f in sorted(os.listdir(path))
33             # if f.endswith(".JPG")
34         ]
35         return np.array([cv2.imread(f) for f in files]), np.array(
36             [cv2.imread(f, 0) for f in files]
37         )
38
39     def non_maximum_suppression(
40         self, image: np.ndarray, window_size: int = 5
41     ) -> np.ndarray:
42         """Apply non-maximum suppression to a given image
43
44         Args:
45             image (np.ndarray): Image to perform non-max suppression on the.
46             window_size (int, optional): The window size around each pixel.
47         Returns:
48             suppressed (np.ndarray): Resulting image after non-maximum suppression.
49         """
50         suppressed = image.copy()
51         global_min = image.min()
52         p = window_size // 2
```

```

53     width, height = suppressed.shape
54
55
56     for i in range(width):
57         x1 = max(0, i - p)
58         x2 = min(width, i + p)
59         for j in range(height):
60             # Bounds for window around pixel at (i, j)
61             y1 = max(0, j - p)
62             y2 = min(height, j + p)
63
64             # Set pixel value to the global min to exclude it from max
65             value = suppressed[i, j]
66             suppressed[i, j] = global_min
67
68             # If pixel has the maximum value in the window then use its value
69             local_max = suppressed[x1:x2, y1:y2].max()
70             if value > local_max:
71                 suppressed[i, j] = value
72
73             # Else keep it is global minimum
74     return suppressed
75
76 def harris_corner_detector(
77     self, image: np.ndarray, k: float = 0.04, window_size: int = 3
78 ) -> np.ndarray:
79     """Given an image, it finds the corners using the Harris Corner Detector.
80
81     Args:
82         image (np.ndarray): Grayscale image to find the corners.
83         k (float, optional): Empirically determined constant in range 0.04 <= x <= 0.06.
84     ↪ Defaults to 0.04.
85         window_size (int, optional): The search window size for operations such as
86     ↪ Gaussian Blur, Sobel Mask etc.. Defaults to 3.
87
88     Returns:
89         corners (list): Corners in the image in the form [(x, y)].
90     """
91
92     # Gaussian Blur over the image
93     image = cv2.GaussianBlur(image, (window_size, window_size), 0)
94
95     # Compute the gradients
96     Ix = cv2.Sobel(image, cv2.CV_64F, dx=0, dy=1, ksize=window_size)
97     Iy = cv2.Sobel(image, cv2.CV_64F, dx=1, dy=0, ksize=window_size)
98
99     # Compute the products of the gradients
100     IxIy = Ix * Iy
101     Ix2 = Ix**2
102     Iy2 = Iy**2
103
104     # Compute the sums of the products of the gradients
105     S2x = cv2.GaussianBlur(Ix2, (window_size, window_size), 0)
106     S2y = cv2.GaussianBlur(Iy2, (window_size, window_size), 0)
107     Sxy = cv2.GaussianBlur(IxIy, (window_size, window_size), 0)
108
109     # Harris Corner Response
110     det = (S2x * S2y) - (Sxy**2)
111     trace = S2x + S2y

```

```

109     R = det - k * (trace**2)
110
111     # Normalize
112     R /= R.max()
113
114     # Non-max suppression
115     R = self.non_maximum_suppresion(R)
116
117     # Thresholding
118     R[R > self.harris_thresh] = 255
119
120     corners = np.where(R == 255)
121     corners = list(zip(corners[0], corners[1]))
122     return corners
123
124 def normalized_cross_correlation(
125     self,
126     image1: np.ndarray,
127     image2: np.ndarray,
128     corners1: list,
129     corners2: list,
130     window_size: int = 5,
131 ) -> dict:
132     """Given two images and their corners, it computes the normalized cross correlation
133     ↪ between the two images and returns the correspondences.
134
135     Args:
136     image1 (np.ndarray): Color image.
137     image2 (np.ndarray): Color image.
138     corners1 (list): It is a list of tuples (x, y) where x and y are the coordinates
139     ↪ of the corners in image1.
140     corners2 (list): It is a list of tuples (x, y) where x and y are the coordinates
141     ↪ of the corners in image2.
142     window_size (int, optional): Window Size. Defaults to 5.
143
144     Returns:
145     Correspondences (dict): Correspondences between the two images in the form
146     ↪ {corner1 (x, y): corner2 (x, y)}
147     """
148     pad = window_size // 2
149
150     # Pad gray images
151     image1_pad = np.pad(image1, pad, "constant", constant_values=0)
152     image2_pad = np.pad(image2, pad, "constant", constant_values=0)
153
154     correspondences = {}
155     for i, corner1 in enumerate(corners1):
156         # Image patch around corner 1
157         x1 = max(corner1[0] - pad, 0)
158         x2 = max(corner1[0] + pad + 1, window_size)
159         y1 = max(corner1[1] - pad, 0)
160         y2 = max(corner1[1] + pad + 1, window_size)
161         patch1 = image1_pad[
162             x1:x2,
163             y1:y2,
164         ]
165
166         max_ncc = -1

```



```

163     best_corner = None
164     for j, corner2 in enumerate(corners2):
165         print(f"i={i + 1}/{len(corners1)} j={j + 1}/{len(corners2)}", end="\r")
166
167         # Image patch around corner 2
168         x1 = max(corner2[0] - pad, 0)
169         x2 = max(corner2[0] + pad + 1, window_size)
170         y1 = max(corner2[1] - pad, 0)
171         y2 = max(corner2[1] + pad + 1, window_size)
172         patch2 = image2_pad[
173             x1:x2,
174             y1:y2,
175         ]
176         # Calculate NCC using image patches
177         patch1_hat = patch1 / np.linalg.norm(patch1)
178         patch2_hat = patch2 / np.linalg.norm(patch2)
179         ncc = np.sum(patch1_hat * patch2_hat)
180
181         # If this NCC is the new max, store it and the coords of the corner
182         if ncc > max_ncc:
183             max_ncc = ncc
184             best_corner = corner2
185
186         # Store correspondence with highest NCC
187         correspondences[corner1] = best_corner
188     return correspondences
189
190 def homography(self, points1: np.ndarray, points2: np.ndarray) -> np.ndarray:
191     """Given two sets of points, it computes the homography matrix.
192
193     Args:
194         points1 (np.ndarray): Points in the form [[x1, y1], [x2, y2], ...] of shape (4 x
195     ↪ 2).
196         points2 (np.ndarray): Points in the form [[x1, y1], [x2, y2], ...] of shape (4 x
197     ↪ 2).
198
199     Returns:
200         h_mat (np.ndarray): Homography matrix of shape (3 x 3).
201     """
202     H = np.zeros((points1.shape[0] * 2, 9))
203     for i, ((x1, y1), (x2, y2)) in enumerate(zip(points1, points2)):
204         # print(f"{i}, ({x1}, {y1}), ({x2}, {y2})")
205         H[2 * i] = [x1, y1, 1, 0, 0, 0, -x2 * x1, -x2 * y1, -x2]
206         H[2 * i + 1] = [0, 0, 0, x1, y1, 1, -y2 * x1, -y2 * y1, -y2]
207     _, _, V = np.linalg.svd(H.T @ H)
208     h = V[-1]
209     h_mat = h.reshape(3, 3)
210     h_mat = h_mat / h_mat[2, 2]
211     return h_mat
212
213 def ransac(
214     self, correspondences: dict, threshold: int = 5, k: int = 72, N: int = 4
215 ) -> np.ndarray:
216     """Performs RANSAC to find the best homography matrix, given the correspondences. This
217     ↪ is done using the largest set of inliers.
218
219     Args:

```

```

217         correspondences (dict): Correspondences between the two images in the form
↪   {corner1 (x, y): corner2 (x, y)}
218         threshold (int, optional): The distance threshold. Defaults to 5.
219         k (int, optional): No of iterations. Defaults to 72.
220         N (int, optional): Sample size. Defaults to 4.
221
222     Returns:
223         H (np.ndarray): Homography matrix.
224         inliers (dict): Inliers in the form {corner1 (x, y): corner2 (x, y)}.
225         outliers (dict): Outliers in the form {corner1 (x, y): corner2 (x, y)}.
226     """
227     H = None
228     inliers = {}
229     outliers = {}
230     max_inliers = 0
231
232     for _ in range(k):
233         set_inliers = {}
234         set_outliers = {}
235
236         num_inliers = 0
237         num_outliers = 0
238
239         # Sample 4 points from the correspondences
240         points1 = choices(list(correspondences.keys()), k=4)
241         points2 = [tuple(correspondences.get(p)) for p in points1]
242
243         # Compute homography matrix using these points
244         h = self.homography(np.asarray(points1), np.asarray(points2))
245
246         for corner1, corner2 in correspondences.items():
247             # Estimate point using homography
248             pt1 = np.array([corner1[0], corner1[1], 1])
249             pt2 = np.array([corner2[0], corner2[1], 1])
250             res = np.dot(h, pt1)
251             # res = np.matmul(h, pt1)
252             res = (res[:2] / res[2]).astype(int)
253             dist = np.linalg.norm(res - corner2)
254
255             # Check if outlier
256             if dist > threshold:
257                 set_outliers[tuple(corner1)] = tuple(corner2)
258                 num_outliers += 1
259                 continue
260
261             # Store inlier
262             set_inliers[tuple(corner1)] = tuple(corner2)
263             num_inliers += 1
264
265         # Check if this homography produced the new largest set of inliers
266         if num_inliers > max_inliers:
267             max_inliers = num_inliers
268             inliers = set_inliers
269             outliers = set_outliers
270             H = h
271
272     return H, inliers, outliers
273

```

```

274 def create_panorama(
275     self, image1: np.ndarray, image2: np.ndarray, H: np.ndarray
276 ) -> np.ndarray:
277     """Creates a panorama image given two images and the homography matrix.
278
279     Args:
280         image1 (np.ndarray): Color image.
281         image2 (np.ndarray): Color image.
282         H (np.ndarray): Homography matrix, H (3x3).
283
284     Returns:
285         Final (np.ndarray): Final panorama image.
286     """
287     copy1, copy2 = image1.copy(), image2.copy()
288     stitcher = cv2.Stitcher().create()
289     final_size = (copy1.shape[1] + copy2.shape[1], copy2.shape[0])
290     final = cv2.warpPerspective(copy2, H, final_size)
291     _, final = stitcher.stitch((copy1, copy2))
292     return final
293
294 def embed_image(
295     self, embed_image: np.ndarray, base_image: np.ndarray
296 ) -> np.ndarray:
297     """Embeds an image into a given base image.
298
299     Args:
300         embed_image (np.ndarray): The image being embedded in the base image.
301         base_image (np.ndarray): The base image being embedded into.
302     Returns:
303         composite_image (np.ndarray): The resulting composite image containing the embed.
304     """
305     # Create a copy of the base image
306     base_image = base_image.copy()
307
308     # Get the desired region to embed image on reom mouse events.
309     region = self.get_user_region(base_image)
310     region = np.asarray(region)
311
312     base_width, base_height, _ = base_image.shape
313     embed_width, embed_height, _ = embed_image.shape
314
315     # Clockwise starting from top left corner
316     embed_image_corners = np.array(
317         [(0, 0), (embed_width, 0), (embed_width, embed_height), (0, embed_height)]
318     )
319
320     # Calculate homography matrix
321     H = self.homography(embed_image_corners, region)
322
323     # Warp the image into the user selected region using the homography matrix
324     embed_image_isolated = cv2.warpPerspective(
325         embed_image, H, (base_height, base_width)
326     )
327
328     # Create a mask of the user selected region
329     fill_mask = np.zeros(base_image.shape).astype("uint8")
330     cv2.fillConvexPoly(fill_mask, region, (255, 255, 255))
331

```

```

332     # Apply mask of user selected region to the base image
333     base_image_with_region = cv2.bitwise_and(base_image, cv2.bitwise_not(fill_mask))
334
335     # Add the base image and the isolated embed image
336     composite_image = base_image_with_region + embed_image_isolated
337
338     return composite_image
339
340 def mouse_callback(
341     self, event: int, x: int, y: int, flags: int, param: dict
342 ) -> None:
343     """Callback function for mouse events used when obtaining the user selected region for
    ↪ embedding an image.
344
345     Args:
346         event (int): The type of mouse click event.
347         x (int): The x pixel coordinate of the mouse in base image.
348         y (int): The y pixel coordinate of the mouse in base image.
349         flags (int): Event flags.
350         param (dict): Any parameters that are passed in.
351     """
352
353     # Handle when the mouse is clicked
354     if event == cv2.EVENT_LBUTTONDOWN:
355         # Get image and corners by reference
356         base_image = param["image"]
357         corners = param["corners"]
358         corners.append([x, y])
359         # Draw square centered around click
360         p = 5
361         start_point = (x - p, y - p) # Top left corner
362         end_point = (x + p, y + p) # Bottom right corner
363         color = (0, 0, 255)
364         thickness = 2
365         cv2.rectangle(base_image, start_point, end_point, color, thickness)
366
367 def get_user_region(self, base_image: np.ndarray) -> list:
368     """Get the corners of the region to embed the image into from the user using mouse
    ↪ click on the base image.
369
370     Args:
371         base_image (np.ndarray): The base image being embedded onto.
372     Returns:
373         corners (list): the pixel coordinates of the corners of the user selected region
    ↪ in the form [(x, y)]
374     """
375     corners = []
376     param = {
377         "image": base_image,
378         "corners": corners,
379     }
380
381     window_name = "Embed Image"
382     cv2.imshow(window_name, base_image)
383     cv2.setMouseCallback(window_name, self.mouse_callback, param)
384
385     while True:
386         cv2.imshow(window_name, base_image)

```

```

387         if cv2.waitKey(1) & 0xFF == ord("q") or len(corners) == 4:
388             cv2.destroyAllWindows(window_name)
389             break
390
391     return corners
392
393 def draw_lines(
394     self, image1: np.ndarray, image2: np.ndarray, points: dict, col1: tuple = None
395 ) -> None:
396     """Draws lines between the points in the two images.
397
398     Args:
399         image1 (np.ndarray): Color image.
400         image2 (np.ndarray): Color image.
401         points (dict): dict of points in the form {corner1 (x, y): corner2 (x, y)}.
402     Return:
403         vis (np.ndarray): Color image with lines.
404     """
405     # Concatenate the two images
406     vis = np.concatenate((image1, image2), axis=1)
407
408     # Draw lines between correlated corners
409     for pt1, pt2 in points.items():
410         col = col1 if col1 else (randint(0, 255), randint(0, 255), randint(0, 255))
411         start_y, start_x = pt1
412         end_y, end_x = pt2
413         end_x = int(end_x + vis.shape[1] / 2)
414         end = (end_x, end_y)
415         start = (start_x, start_y)
416         cv2.line(vis, start, end, col, 1)
417
418     return vis
419
420 def draw_corners(self, image: np.ndarray, corners: list) -> None:
421     """Draw corners in the image.
422
423     Args:
424         image (np.ndarray): Color image.
425         corners (list): List of points in the form [(x, y)].
426     Returns:
427         vis (np.ndarray): Color image with corners drawn.
428     """
429     vis = image.copy()
430     # Draw lines between correlated corners
431     for corner in corners:
432         cv2.circle(vis, (corner[1], corner[0]), 3, (0, 0, 255), 1)
433     return vis
434
435 def show_image(self, images: list, titles: list) -> None:
436     """Shows a list of images along with their titles.
437
438     Args:
439         images (list): A list of all the images to be shown.
440         titles (list): A list of all the titles of the images.
441     """
442     for image, title in zip(images, titles):
443         cv2.imshow(title, image)
444     if cv2.waitKey(0) & 0xFF == ord("q"):

```

```

445         cv2.destroyAllWindows()
446
447     def save_image(self, images: list, fnames: list) -> None:
448         """Saves a list of images with their corresponding filename
449
450         Args:
451             images (list): A list of all images to be saved
452             fnames (list): A list of all the file names of the images
453         """
454         for image, fname in zip(images, fnames):
455             cv2.imwrite(f"results/{fname}.jpg", image)
456
457
458     def main():
459         """Main function to run the Panorama class."""
460         # The directory to use two images from to create a mosaic
461         DIR = "DanaHallWay1"
462         # DIR = "DanaOffice"
463         pano = Panorama(DIR)
464
465         # Load images
466         col, gray = pano.load_images(DIR)
467         image1, image2 = col[0], col[1]
468         gray1, gray2 = gray[0], gray[1]
469
470         # Apply Harris corner detector
471         corners1 = pano.harris_corner_detector(gray1)
472         corners2 = pano.harris_corner_detector(gray2)
473
474         # Find correspondences using NCC
475         correspondences = pano.normalized_cross_correlation(
476             gray1, gray2, corners1, corners2
477         )
478
479         # Use RANSAC to estimate homography matrix and find inliers
480         H, inliers, outliers = pano.ransac(correspondences)
481
482         # Warp images
483         output = pano.create_panorama(image1, image2, H)
484
485         # Create visualiations of results
486         corners1_vis = pano.draw_corners(image1, corners1)
487         corners2_vis = pano.draw_corners(image2, corners2)
488         correspondences_vis = pano.draw_lines(image1, image2, correspondences)
489         inliers_vis = pano.draw_lines(image1, image2, inliers)
490         outliers_vis = pano.draw_lines(image1, image2, outliers)
491
492         # Display results
493         print(H)
494         pano.show_image(
495             [
496                 image1,
497                 image2,
498                 corners1_vis,
499                 corners2_vis,
500                 correspondences_vis,
501                 inliers_vis,
502                 outliers_vis,

```

```

503         output,
504     ],
505     [
506         "Input 1",
507         "Input 2",
508         "corners1",
509         "corners2",
510         "correspondences",
511         "inliers",
512         "outliers",
513         "Output",
514     ],
515 )
516
517 # Save results
518 pano.save_image(
519     [
520         image1,
521         image2,
522         corners1_vis,
523         corners2_vis,
524         correspondences_vis,
525         inliers_vis,
526         outliers_vis,
527         output,
528     ],
529     [
530         f"{DIR}_input1",
531         f"{DIR}_input2",
532         f"{DIR}_corners1",
533         f"{DIR}_corners2",
534         f"{DIR}_correspondences",
535         f"{DIR}_inliers",
536         f"{DIR}_outliers",
537         f"{DIR}_output",
538     ],
539 )
540
541
542 def extra_credit():
543     """Extra credit function that embeds an image into another image."""
544     # The directory to pull an image from to use as the base image
545     DIR = "DanaHallWay1"
546     # DIR = "DanaOffice"
547
548     pano = Panorama(DIR)
549
550     # Load images
551     col, _ = pano.load_images(DIR)
552     embed_images, _ = pano.load_images("ec")
553
554     embed_image = embed_images[0]
555     # embed_image2 = embed_images[1]
556     base_image = col[0]
557
558     # Warp an image into a region in the second image
559     output = pano.embed_image(embed_image, base_image)
560

```

```
561     # Display results
562     pano.show_image([output], ["ExtraCredit_output"])
563
564     # Save Results
565     pano.save_image([output], ["ExtraCredit_output"])
566
567
568 if __name__ == "__main__":
569     main()
570     extra_credit()
```